

Recommender Systems

推薦系統

賴旭昭 Hsu-Chao Lai

National Cheng Kung University



About Me

- Hsu-Chao Lai 賴旭昭
- AI Scientist at NetDB Lab
- Ph.D. of NYCU CS
- Specialized in
 - Recommender systems
 - Real-time bidding
 - AI stock trading
- hclai@netdb.csie.ncku.edu.tw

Outline

- Introduction
- Content-based Approaches
- Collaborative Filtering
- Matrix Factorization
- Ranking Objective
- Takeaways
- Appendix: Glossary of Hot Topics
- References

Background Knowledge:

- Linear regression
- Logistic regression
- Maximum log-likelihood
- Or Minimum negative log-likelihood
- Singular Value Decomposition (SVD)
- (Stochastic) gradient descent

Example

momo



别人也逛過

熱銷排行

1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10

今日 下單登記抽! NEBULA 無線微型投影機!

電玩 指品滿額 搶2000 MO幣

【Nintendo 任天堂】Switch 超級瑪利歐兄弟 驚奇(中文版)

\$1,290

今日 下單登記抽! NEBULA 無線微型投影機!

電玩 指品滿額 搶2000 MO幣

【Nintendo 任天堂】Switch OLED 白色主機+健身環大冒險同捆組+遊戲選一

折價券 \$13,660

今日 下單登記抽! NEBULA 無線微型投影機!

電玩 指品滿額 搶2000 MO幣

【Nintendo 任天堂】Switch 紅藍主機 電力加強版 日規+遊戲選一

折價券 \$9,990

今日 下單登記抽! NEBULA 無線微型投影機!

電玩 指品滿額 搶2000 MO幣

【Nintendo 任天堂】Switch 瑪利歐明星遊戲系列多選一(台灣公司)

\$1,283

1 | 2 | 3 | 4 | 5 | 6 | 7 | 8

今日 下單登記抽! NEBULA 無線微型投影機!

電玩 指品滿額 搶2000 MO幣

【SONY 索尼】New PlayStation 5 光碟版主機(PS5 Slim)(CFI-70081)

\$17,580

今日 下單登記抽! NEBULA 無線微型投影機!

電玩 指品滿額 搶2000 MO幣

【SONY 索尼】New PlayStation 5 數位版主機(PS5 Slim)(CFI-70082)

\$14,580

今日 下單登記抽! NEBULA 無線微型投影機!

電玩 指品滿額 搶2000 MO幣

【SONY 索尼】PS5 Pro 遊戲主機 - PlayStation 5 Pro(CFI-7022B01)


\$24,280

今日 下單登記抽! NEBULA 無線微型投影機!

電玩 指品滿額 搶2000 MO幣

【Nintendo 任天堂】Switch電續加強主機+健身環同捆組+專用豪華收

折價券 \$9,990



► Types of Recommendations

- Non-personalized recommendations
 - Simple aggregation
 - Top-10 popular, new-on-shelf, most advertised
- Personalized recommendations
 - Specially tailored to **each individual**
 - Find relevant items that are now browsing
 - Amazon, Netflix, YouTube, Facebook, news, etc
 - Customize user experience
 - Spotify, trip planning, friend, etc
 - Increase conversion rate (e.g., purchase)
 - Ads, new items, etc

Recommender Systems (RS)

- U : a set of **users**
- V : a set of **items**
- Optimize a **Utility function** $f: U \times V \rightarrow R$
 - R : a set of **user feedback**
 - Ratings (1-5 stars)
 - Binary (purchased or not)
 - Polarity (like or dislike)
 - Preference, satisfaction, etc

你喜歡這篇文章嗎?
還沒有人評價此篇文章



Din Tai Fung 101

鼎泰豐 101店

4.5 ★★★★★ (11,127) • \$400–800
Taiwanese restaurant • ♂

Workflow

1. Collect **known user feedback**

- What data can be used as user feedback?

2. Infer **unknown user feedback** based on the known ones

- Focus on retrieving unknown feedback with high ratings
 - Low rating ones do not bring profits

3. Evaluate the inference of **unknown ones**

- How to properly measure the performance of recommendations

► (1) Known User Feedback Collection

- **Explicit Feedback**
 - Ask users to rate items
 - 1-5 stars
 - Very dissatisfied, dissatisfied, moderate, satisfied, very satisfied
 - Crowdsourcing: pay users to rate
- **Implicit Feedback**
 - User actions/engagement
 - *Purchase* implies high ratings
 - *Watched* implies high ratings
 - Watch time, visit and re-visit, etc
 - What defines low ratings?

Utility Matrix with Explicit Feedback


Movie Rating Records		
Willie	全面啟動	5
Xavier	星際效應	2
Zack	全面啟動	4
Willie	星際效應	5
Xavier	鬼滅之刃	5
Willie	奧本海默	5
Zack	奧本海默	4



Ratings	奧本海默	星際效應	全面啟動	鬼滅之刃	天氣之子
Willie	5	5	5		
Xavier		2		5	
Yvonne					
Zack	4		4		

Utility Matrix with Implicit Feedback


Movie Watching Records		
Willie	全面啟動	Watched (5)
Xavier	星際效應	Watched (2)
Zack	全面啟動	Watched (4)
Yvonne	天氣之子	Watched (-)
Willie	星際效應	Watched (5)
Xavier	鬼滅之刃	Watched (5)
Yvonne	鬼滅之刃	Watched (-)
Willie	奧本海默	Watched (5)
Zack	奧本海默	Watched (4)
Zack	星際效應	Watched (-)



	Watched	奧本海默	星際效應	全面啟動	鬼滅之刃	天氣之子
Willie	1	1	1			
Xavier			1		1	
Yvonne					1	1
Zack	1	1	1	1		

► (2) Inferring Utilities

- Key challenges:
 - Utility matrix M is **sparse**
 - Most of the entries of M are left blank
 - Sparsity¹: MovieLens 95-99%, Netflix Prize 99%, Amazon Product Review 99.9%
 - **Cold-start** issues
 - New items or users have few records to learn
- Three classic approaches:
 - **Content-based**
 - **Collaborative Filtering**
 - **Latent-based**



¹[MovieLens](#) | [GroupLens](#), [Netflix Prize data](#), and [Amazon Review Datasets](#)

► (3) Evaluation

- Pointwise objectives
 - Classification: logistic function, cross-entropy, softmax, ...
 - Regression: MSE, MAE, Mean Squared Logarithmic Error, ...
- **Highly sensitive** to the **sparsity** issue
- Given the ratio of positive to negative data is about 1:99 in M
- Always inferring “negative” yields 99% accuracy during training
- Unable to retrieve items with high ratings
- Solution: **pairwise objectives (a.k.a. ranking)**



Outline


- Introduction
- Content-based Approaches
- Collaborative Filtering
- Matrix Factorization
- Ranking Objective
- Takeaways
- Appendix: Glossary of Hot Topics
- References

Content-based Approaches

- Core idea: **items have profiles!**
 - Movie: {genre, director, casts, series}
 - Article: {author, platform, a set of keywords}
 - E-commerce product: {function, brand, specs}
- Recommend unknown items to customer u similar to previous items highly rated by u
- Recommend C. Nolan's new movie to Willie
- Recommend new Apple products to loyal customers

Watched	奧本海默	星際效應	全面啟動
Willie	1	1	1

Method Plan



► (2) Feature Extraction for Item Profiles

- A set (vector) of features *distinguishing each item*
- Structural data
 - Movie recommendation: {genre, director, casts, series}
- Unstructured data
 - **Text:** **key words**, latent topics, writing style, ...
 - News, product descriptions, reviews, blogs, social media, etc
 - **Video:** latent topics, sequential storytelling, ...
 - **Image:** objects, styles, color tones, ...
 - *Too abstract to extract*

A Glimpse of Key Word Extraction: TF-IDF

1. Frequency $f_{x,y}$ of word x in document y
 - $TF_{x,y} = f_{x,y} / \text{MAX}_{\forall i} f_{i,y}$
2. The **uniqueness** of x among all documents
 - $IDF_{x,y} = \ln(N/n_x)$,
 - where N is the number of total documents, and n_x is the number of docs containing x
3. TF-IDF score: $w_{x,y} = TF_{x,y} \times IDF_{x,y}$
4. **Document profile:** a set of words with top- k TF-IDF scores

Modern approach key words: text embeddings, LLM, attention, transformer, BERT, ...

► (3)(4) User Profiles and Recommendations

- A set (vector) of features *profiling/describing user preference*
 - Weighted average of rated item profiles
- Matching user profile \mathbf{u} to unknown item profile \mathbf{v}
- **Cosine similarity:** $f(\mathbf{u}, \mathbf{v}) = \cos(\mathbf{u}, \mathbf{v}) = \frac{\mathbf{u} \cdot \mathbf{v}}{\|\mathbf{u}\| \cdot \|\mathbf{v}\|}$
- Recommend the unknown items with *top-k cosine similarity scores*

▶ Pros and Cons for Content-based Methods

Pros	Cons
No sparsity issue	Extensive feature engineering
No cold-start item problem	Cold-start user problem
Identify unique items or tastes of users	Poor at exploring new items
Interpretable	



Outline

- Introduction
- Content-based Approaches
- Collaborative Filtering
- Matrix Factorization
- Ranking Objective
- Takeaways
- Appendix: Glossary of Hot Topics
- References

What If Content Is Unavailable

- And sometimes engineering features is just too expensive

Watched	奧本海默	星際效應	全面啟動	鬼滅之刃	天氣之子
Willie	5	5	5		
Xavier				5	
Yvonne					
Zack	4		4		

What If Content Is Unavailable

Given records in blue

Watched	奧本海默	星際效應	全面啟動	鬼滅之刃	天氣之子
Willie	5	5	5		
Xavier		?		5	
Yvonne		?			
Zack	4	?	4		

Who is likely to watch this movie? Why?

Collaborative Filtering (CF)

- [Assumption] **Similar users** (based on known) tend to have **similar behavior/ratings** (on unknown)

⇒ User-user Collaborative Filtering

⇒ Item-item Collaborative Filtering


- No need of content
- Treat **behavior** as features

- **Similarity measurement** is the key question of CF!

● 別人也逛過



Method Plan



► (1) Choices of Similarity Function

- Cosine Similarity (numerical, binary)
 - Euclidean distance (numerical)
 - Jaccard similarity (binary)
 - Pearson correlation (numerical)
-
- Choose function based on data types
 - Be careful of the limitations and cons

Jaccard Similarity

u_i : user

I_i : interacted items of u_i

- For users u_i and u_j , $Jacc(i, j) = \frac{|I_i \cap I_j|}{|I_i \cup I_j|} \in [0, 1]$
 - 0 means no overlap, while 1 means totally the same
- $I_i = \{a, b, c, d, e\}$ and $I_j = \{a, b, c, x, y\}$
- $Jacc(i, j) = 3/7$
- Cons 1: when item space is large, jaccard is easily diluted to 0
 - Every pair of user unions exactly 100,000,000 items
- Cons 2: unable to handle numerical values and negative records
- And there are more ...

Pearson Correlation

i, j : user identifier

I_i : interacted items of user i

$R_{i,v}$: ratings for item v of user i

\bar{R}_i : average ratings of user i

$$\bullet P_{i,j} = \frac{\sum_{v \in I_i \cap I_j} (R_{i,v} - \bar{R}_i)(R_{j,v} - \bar{R}_j)}{\sqrt{\sum_{v \in I_i \cap I_j} (R_{i,v} - \bar{R}_i)^2} \sqrt{\sum_{v \in I_i \cap I_j} (R_{j,v} - \bar{R}_j)^2}}$$

- Recall cosine similarity $\cos(u, v) = \frac{u \cdot v}{\|u\| \cdot \|v\|}$
- **Normalized variation** of cosine similarity
- Cons: inaccurate when intersection set is small
 - Every pair of user shares exactly one interacted item

► (2)(3) Rating Predictions

- Let S_i^U be the set of top- k similar users of user i
- Let $r_{i,v}$ be the ratings of item v for user i
- **Average score:** $r_{i,v} = \frac{1}{k} \sum_{j \in S_i^U} r_{j,v}$
- **Similarity-weighted score:** $r_{i,v} = \frac{\sum_{j \in S_i^U} \text{sim}(i,j) \times r_{j,v}}{\sum_{j \in S_i^U} \text{sim}(i,j)}$
- **kNN-classifier:** output the majority class/rating of S_i^U

Item-item CF

GUESS: In practice, which is better?
User-user CF or item-item CF?

- Similar items may attract similar user interactions
- Apply same tricks as user-user CF

$$r_{u,i} = b_{u,i} + \frac{\sum_{j \in S_i^I} sim(i,j) \times (r_{u,j} - b_{u,j})}{\sum_{j \in S_i^I} sim(i,j)}$$

- S_i^I is the set of top- N similar items
- $sim(i,j)$ here is the item-item similarities

- $b_{u,i} = \mu + b_u + b_i$ is the baseline estimation of $r_{u,i}$
- μ is the overall mean item ratings
- b_u is the rating deviation (bias) of user $u \Rightarrow$ (avg. ratings of u) - μ
- b_i is the rating deviation (bias) of item i

▶ Pros and Cons for Collaborative Filtering

Pros	Cons
No need of feature engineering	Cold-start item/user problem
	Sparsity issue
	Tends to overlook unpopular items

Is it possible to combine content-based and CF to enjoy both benefits? How?

Outline

- Introduction
- Content-based Approaches
- Collaborative Filtering
- Matrix Factorization
- Ranking Objective
- Takeaways
- Appendix: Glossary of Hot Topics
- References

Review CF

- $$r_{u,i} = b_{u,i} + \frac{\sum_{j \in S_i^I} sim(i,j) \times (r_{u,j} - b_{u,j})}{\sum_{j \in S_i^I} sim(i,j)}$$
- Is similarity-weighted average flexible?
 - Strong and **artificial** assumption restricts learning
 - May overlook **interdependencies** between users
- How to design **a similar but data-driven approach?**
 - Hint: it's a weighted average formulation

- S_i^I is the set of top- N similar items
- $sim(i,j)$ here is the item-item similarities
- $b_{u,i} = \mu + b_u + b_i$ is the baseline estimation of $r_{u,i}$
- μ is the overall mean item ratings
- b_u is the rating deviation of user u
- b_i is the rating deviation of item i

Weighted Sum Approach

- Use **weighted sum** rather than weighted average!

$$\widehat{r}_{u,i} = b_{u,i} + \sum_{j \in S_{u,i}^I} w_{i,j} (r_{u,j} - b_{u,j})$$

- $S_{u,i}^I$: set of items rated by user u and similar to item i
- $w_{i,j}$: weights between similar item pairs
- How to obtain $w_{i,j}$?
- Regression! **Minimize SSE**: $\sum_{(u,i) \in R} (\widehat{r}_{u,i} - r_{u,i})^2$

R contains every rated pair of user u and item i

Optimization

- Objective function SSE:

$$L(\mathbf{w}) = \sum_{(u,i) \in R} \frac{([b_{u,i} + \sum_{j \in S_{u,i}^I} w_{i,j} (r_{u,j} - b_{u,j})] - r_{u,i})^2}{\text{Predicted rating}} \quad \text{True rating}$$

- Gradient Descent!
- $\mathbf{w} \leftarrow \mathbf{w} - \eta \nabla_{\mathbf{w}} L$
- Still treat users and items **independently**.
 - $f(u, v) = f(u) + f(v)$

Modeling User-item Interaction

- Linear formulation: $\widehat{r}_{u,i} = \mu + b_u + b_i$
 - μ is the overall mean item ratings
 - b_u is the user bias
 - b_i is the item bias
- Second-order: $\widehat{r}_{u,i} = \mu + b_u + b_i + \sum_{u,i} (w_u x_u) \times (w_i x_i)$
 - Optimizing w_u and w_i is affected by $x_u x_i \leq \text{interaction!}$
 - x_u, x_i are user and item features
 - w_u, w_i are weights
- Expensive features
- High dimensionality of features might overwhelm learning
- Few samples to learn well due to sparsity

Matrix Factorization (MF)

- Reduce to: $\widehat{r}_{u,i} = \mu + b_u + b_i + \sum_{u,i} p_u q_i$

p_u, q_i are some **easy-to-get**, and **low-rank** features of user and item

- If we ignore μ, b_u, b_i and rewrite to matrix form: $R \approx P Q^T$

R is the utility matrix

- $P \in \mathbb{R}^{|U| \times m}$ compose latent features of users
- $Q \in \mathbb{R}^{|V| \times m}$ compose latent features of items

3				3	3		1		3
1			3	1		1	5		
	3				1	4		5	4
	1	1			3	5			3
2		1		5			3		
R	3	1	1	4	4		5		

\approx

8.0	-1.5	-0.1
7.2	2.9	1.6
10.	-1.7	2.4
8.5	-1.0	-1.6
10.	2.3	-1.3
8.5	-0.7	-1.0

.28	.24	.21	.35	.33	.31	.35	.36	.34	.33
-.35	.15	-.12	.06	-.25	.34	-.34	.70	-.16	-.11
-.03	.26	.53	.00	-.22	-.61	-.29	.17	.31	.08

Q^T

P

6 users vs 10 items, and $m = 3$

Is MF What We Want?

R

3				3	3		1		3
1			3	1		1	5		
	3				1	4		5	4
	1	1			3	5			3
2		1			5			3	
3	1	1	4	4			5		

\approx

8.0	-1.5	-0.1
7.2	2.9	1.6
10.	-1.7	2.4
8.5	-1.0	-1.6
10.	2.3	-1.3
8.5	-0.7	-1.0


.28	.24	.21	.35	.33	.31	.35	.36	.34	.33
-.35	.15	-.12	.06	-.25	.34	-.34	.70	-.16	-.11
-.03	.26	.53	.00	-.22	-.61	-.29	.17	.31	.08

P

Q^T

- Easy-to-get data? **Yes!**
 - Need only the utility matrix
 - No additional feature engineering
- Low-rank? **Yes!**
 - On-demand size m
- Wait, what is this exactly?

Latent Factor (a.k.a. Embedding)



- Latent factor is the **hidden feature** that captures underlying patterns in the behavioral data
 - Abstractive concepts of directors, genres, and popularity
 - Data-driven
- Interpretability ... Meh ...
- Optimization? Hint: linear algebra

First Approach: SVD

- Matrix factorization $R \approx PQ^T$
- **Singular Value Decomposition** (SVD): $A \approx U\Sigma V^T$
 - SVD always yields minimal reconstruction loss: $\sum_{i,j} (A_{i,j} - [U\Sigma V^T]_{i,j})^2$
 - Wait... isn't it SSE!? JACKPOT!
- **Sadly, SVD isn't designed for missing values**
- Need other approaches to optimize MF

$$A \approx U \Sigma V^T$$

Second Approach: Gradient Descent (GD)

- Revisit objective function:

$$\min_{\mathbf{P}, \mathbf{Q}} L^{MF} = \sum_{(u,i) \in R} (r_{u,i} - \hat{r}_{u,i})^2 = \sum_{(u,i) \in R} (r_{u,i} - p_u \cdot q_i)^2$$

- To prevent overfitting, we add L2 regularization:



$$\min_{\mathbf{P}, \mathbf{Q}} L^{MF} = \sum_{(u,i) \in R} (r_{u,i} - p_u \cdot q_i)^2 + \lambda_1 \sum_u \|p_u\|^2 + \lambda_2 \sum_i \|q_i\|^2$$

- Initialize \mathbf{P}, \mathbf{Q} with small non-zero random values
- Do gradient descent:
 - $\mathbf{P} \leftarrow \mathbf{P} - \eta \nabla_{\mathbf{P}} L^{MF}$
 - $\mathbf{Q} \leftarrow \mathbf{Q} - \eta \nabla_{\mathbf{Q}} L^{MF}$

Stochastic Gradient Descent (SGD)

- Matrix operation was **slow** for CPU
(but fast for GPU)
- **GD:** $P \leftarrow P - \eta \sum_{r_{u,i}} \nabla_P P(r_{u,i})$
- **SGD:** $P \leftarrow P - \eta \nabla_P P(r_{u,i})$
- **Evaluate on all ratings** vs.
- **Evaluate on one rating**
- **Fast convergence**

SGD vs. GD



GD improves the value of the objective function at every step.
SGD improves the value but in a “noisy” way.
GD takes fewer steps to converge but each step takes much longer to compute.
In practice, **SGD** is much faster!

Add Biases Back

- Overall objective function

$$\min_{\mathbf{P}, \mathbf{Q}, \mathbf{b}} L^{MF} = \sum_{(u,i) \in R} [r_{u,i} - (\mu + b_u + b_i + p_u \cdot q_i)]^2$$

Goodness of fit

$$+ \lambda_1 \sum_u \|p_u\|^2 + \lambda_2 \sum_i \|q_i\|^2 + \lambda_3 \sum_u \|b_u\|^2 + \lambda_4 \sum_i \|b_i\|^2$$

Regularization

- Learn p_u, q_i, b_u, b_i with SGD
- Select λ with the validation set
- Further reading: Alternating Least Square (ALS)

▶ Pros and Cons for Matrix Factorization

Pros	Cons
No need of feature engineering	Poor interpretability
Model user-item interactions	Additional hyperparameters λ, m
Capture hidden relationships	Cold-start user and item problem
Data-driven	Sparsity issue



Outline

- Introduction
- Content-based Approaches
- Collaborative Filtering
- Matrix Factorization
- Ranking Objective
- Takeaways
- Appendix: Glossary of Hot Topics
- References

Pointwise Recommendations

- Classification, regression, etc
 - “Predicting whether a user will rate a movie 5 stars.”
 - “Predicting the rating a user will rate.”
 - **Use rated data only =>** abandon huge amount of unrated data
 - **Implicit feedback =>** severe imbalance; unknown or dislike?
- **Rethink: is rating that important in recommendations?**
 - Given movie prediction={A:4.6, B:4.7, C:4.5}, recommend which one?
- **Feasible to implicit feedback**
- **Controllable imbalance**

Relative Preference

- Another mindset: **Which one is better? Item i or j ?**
- We care about their **relative preference**, not their ratings
 - $i > j$ or $i < j$?
 - Generalize definition of preference <- **implicit feedback**
- Given user u likes item i better than j
- For u , maximize the **probability of ranking item i over j**
 - Pointwise: $\min \sum (\widehat{r}_{u,i} - r_{u,i})^2$
 - Pairwise: **$\max \sum P(\widehat{r}_{u,i} > \widehat{r}_{u,j} \mid r_{u,i} > r_{u,j})$**

Purchase > add-to-cart
Click > seen
Finish > skip
Interacted > not interacted

Pairwise Relation Dataset

- Define $i >_u j$, where u interacted with i but not j
- A training database O collects triplets (u, i, j) s.t. $i >_u j$
 - Negative sample ratio? Let's say we collect just enough to train for now

	i_1	i_2	i_3	i_4
u_1	1	1		
u_2		1		1
u_3				1
u_4	1			1

$$O = \left\{ (u_1, i_1, i_3), (u_1, i_2, i_3), (u_1, i_1, i_4) \right\}$$

Bayesian Personalized Ranking (BPR)

- Let Θ denote be learnable model parameter, we maximize

$$\max \sum P(\hat{r}_{u,i} > \hat{r}_{u,j} \mid r_{u,i} > r_{u,j}) \quad \frac{p(\Theta | i >_u j)}{\text{Likelihood!}} \times \frac{p(\Theta)}{\text{Follow Gaussian Distribution}}$$


- Recall **sigmoid function** (**maximum likelihood** for logistic regression)

$$p(i >_u j | \Theta) = \sigma(\hat{r}_{u,i,j}(\Theta)) = \sigma(\hat{r}_{u,i,j})$$

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

BPR Objective

- Maximum log-likelihood function
- $BPR - Opt = \prod_{(u,i,j) \in O} p(\Theta | i >_u j)$
- $= \ln \prod_{(u,i,j) \in O} p(i >_u j | \Theta) p(\Theta)$
- $= \sum_{u,i,j} \ln \sigma(\hat{r}_{u,i,j}) + \ln p(\Theta)$
- $= \sum_{u,i,j} \ln \sigma(\hat{r}_{u,i,j}) - \frac{\lambda}{2} \|\Theta\|^2$
- $\hat{r}_{u,i,j}$ should quantify $i >_u j$
 - $\hat{r}_{u,i,j} = \hat{r}_{u,i} - \hat{r}_{u,j} = \begin{cases} > 0 & \text{if ranked correctly} \\ \leq 0 & \text{else} \end{cases}$



SGD



- $\Theta \leftarrow \Theta - \eta \nabla_{\Theta}$
- $\nabla_{\Theta} = \frac{\partial \text{BPR}-Opt}{\partial \Theta} = \dots = \sum_{u,i,j} \frac{e^{-(\hat{r}_{u,i} - \hat{r}_{u,j})}}{1 + e^{-(\hat{r}_{u,i} - \hat{r}_{u,j})}} \cdot \frac{\partial}{\partial \Theta} (\hat{r}_{u,i} - \hat{r}_{u,j}) - \lambda \Theta$
- Fix $\hat{r}_{u,j}$ and learn $\hat{r}_{u,i}$, and vice versa
- MF-based formulation: $\hat{r}_{u,i} = \mu + b_u + b_i + p_u \cdot q_i$
- What is different for MF now?

BPR-MF vs. Vanilla MF

	BPR-MF	Vanilla MF (SVD)
Objective	Pairwise	Pointwise
Latent Factor	Explain relative orders	Explain ratings
Gradient of User	$f_u(x_i - x_j)$ based on feature deviation	$f_u(x_i)$ based on item feature
Gradient of Positive Item	$f_i(x_u)$	$f_i(x_u)$
Gradient of Negative Item	$f_j(-x_u)$ shift away from user feature	None


Performance Evaluation (AUC)

- AUC is the probability that a randomly chosen positive item is **ranked higher** than a randomly chosen negative item




Pros and Cons of BPR

Pros	Cons
Directly model relative preference	Ignore absolute preferences
Incorporate implicit feedback	Computationally-intensive
Handles sparsity well	Cold-start problem
Great compatibility	Additional hyperparameter: negative sample rate




Takeaways

- Choose models according to data and scenarios
- Combine any of them is possible
- Cold-start and sparsity issue
- Compatible to new tech (e.g., NN)



Agentic Recommender System Era

- Pros
 - Easy to integrate multimodality
 - Unlimited feedback augmentation
 - Dynamic and continuous self-improvement
 - Flexible to extend novelty
- Cons
 - Intensive computational cost
 - Unintended outcomes
 - Unstable performances
 - Overamplified biases (data poisoning)
 - Security and privacy concerns



Appendix: Glossary of Hot Topics

- Adversarial Training
- Attention Network
- Contrastive Learning
- Cross-domain Rec.
- Data Augmentation
- Fairness Rec.
- Federated Learning
- Generative Models
- Graph Neural Networks
- Knowledge Graph-based Rec.
- Meta-learning
- Multi-modal Learning
- Neural Collaborative Filtering
- Self-supervised Learning
- Sequence Rec.
- Session-based Rec.
- Social, spatial, and temporal Rec.
- Transformer-based Rec.

► References and Further Reading

- Jure Leskovec, *Stanford CS246 lecture note*, 2024
- Jingbo Shan, *UCSD DSC148 lecture note*, 2023
- Steffen Rendle et al.: BPR: *Bayesian Personalized Ranking from Implicit Feedback*. UAI 2009